


Lab 2: Unsupervised Learning

K-Means, Gaussian Mixture Models (GMM), Expectation Maximization (EM)

Machine Learning and Deep Learning 2023

Eros Fani & Debora Caldarola

eros.fani@polito.it, debora.caldarola@polito.it





Unsupervised learning

What is unsupervised learning?

- Introduction to clustering
 - Introduction
 - Metrics
- K-Means
 - how does it work?
 - PROs, CONs
- GMMs
 - what is a GMM?
 - How does it differ from K-means?
 - PROs, CONs



Unsupervised Learning

- **Goal:** discover patterns or structure within a dataset without the use of explicit labels or guidance
 - Identify the hidden structure behind the data
- No feedback mechanism for the algorithm to determine how close or far it is from the optimal solution




Unsupervised Learning

- Input dataset $X = \{x_1, x_2, \dots, x_n\}$, $x_i \in \mathbb{R}^m$
- **Goal:** find $f: X \rightarrow Y$ mapping the input data to a set of output variables Y
 - f is not provided with explicit target variables or labels
 - f should discover underlying patterns, structure or relationships within the data
 - f should be able to generalize to new, unseen data
- Particularly useful for problems where the desired outcome is not known in advance
- **EXAMPLES**
 - clustering
 - dimensionality reduction
 - anomaly detection

Clustering

- Unsupervised ML technique involving grouping similar data points together based on their similarity or distance from each other
- **Goal:** identify natural groupings, or **clusters**, in a dataset without any prior knowledge or labels
- Good clusters if:
 - high intra-cluster similarity
 - low inter-cluster similarity
- **EXAMPLES:**
 - **k-means clustering**
 - hierarchical clustering
 - DBSCAN clustering
 - density-based clustering (GMM)





Clustering metrics

SUPERVISED CLUSTERING METRICS:

- Require true labels to be computed
- Examples: homogeneity, completeness, V-measure

UNSUPERVISED CLUSTERING METRICS:

- Do not require true labels to be computed
- Example: Silhouette coefficient

<https://scikit-learn.org/stable/modules/clustering.html#homogeneity-completeness>

<https://scikit-learn.org/stable/modules/clustering.html#silhouette-coefficient>

K-Means


- 1) Sample K centroids randomly
- 2) Assign each sample to the nearest centroid
- 3) Compute new cluster centroids
- 4) Repeat until convergence

PROs:

- Simple and efficient
- Scalable
- Interpretable results
- Flexible
- Guaranteed to converge

CONs:

- Sensible to initialization
- Difficult to choose K
- Sensitive to outliers
- Limited to linear boundaries
- Biased towards equal-sized clusters
- Only handles numerical data
- Assumes spherical clusters



Exercise 1: K-Means with Sklearn

Explore the K-means algorithm on the Iris dataset

- Import the Iris dataset
- Evaluate K-means with 3 clusters (= 3 real labels: Virginica, Versicolour, Setosa) with several random seeds
 - How does the random initialization of the centroids affect the final clusters?
- Evaluate K-means with a fixed random seed, with $K = 2, 3, \dots, 20$
 - What is the best K according to the supervised clustering metrics? And according to the unsupervised clustering metrics? Why?
- Reduce the data with PCA and apply k-means again. Visualize the results using PCA with 2 components (suggestion: use the script in the next slides)



How to visualize the results (1/3)

```
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
import numpy as np

reduced_data = PCA(n_components=2).fit_transform(iris.data)
kmeans_model.fit(reduced_data)

# Step size of the mesh. Decrease to increase the quality of the VQ.
h = 0.02 # point in the mesh [x_min, x_max][y_min, y_max].

# Plot the decision boundary. For that, we will assign a color to each
x_min, x_max = reduced_data[:,0].min() - 1, reduced_data[:,0].max() + 1
y_min, y_max = reduced_data[:,1].min() - 1, reduced_data[:,1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# Obtain labels for each point in mesh. Use last trained model.
Z = kmeans_model.predict(np.c_[xx.ravel(), yy.ravel()])
```

How to visualize the results (2/3)

```
# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.figure(1)
plt.clf()
plt.imshow(
    Z,
    interpolation="nearest",
    extent=(xx.min(), xx.max(), yy.min(), yy.max()),
    cmap='Set2',
    aspect="auto",
    origin="lower",
)

plt.scatter(reduced_data[:,0], reduced_data[:,1], c=iris.target, cmap=Pastel12)
# Plot the centroids as a white X
centroids = kmeans_model.cluster_centers_
```

How to visualize the results (3/3)

```
plt.scatter(
    centroids[:, 0],
    centroids[:, 1],
    marker="x",
    s=169,
    linewidths=3,
    zorder=10,
)
plt.title(
    "K-means clustering on the digits dataset (PCA-reduced data)\n"
    "Centroids are marked with white cross"
)
plt.xlim(x_min, x_max)
plt.ylim(y_min, y_max)
plt.xticks(())
plt.yticks(())
plt.show()
```

K-means clustering on the digits dataset (PCA-reduced data)
Centroids are marked with white cross




Image from sklearn official tutorials



Tips

- Remember to standardize the data
- Refer to the following documentation:
 - https://scikit-learn.org/stable/auto_examples/datasets/plot_iris_dataset.html
 - <https://scikit-learn.org/stable/modules/clustering.html#silhouette-coefficient>
 - <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>
 - <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>
 - <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>




Gaussian Mixture Models (GMMs)

- Overcoming the drawbacks of K-means
- The idea behind GMMs
- GMMs for handling not well separable data
- GMMs for generating new data



Main issues with *K*-means


- Poor performance in real-world scenarios due to the distance-from-cluster idea used to determine cluster membership






Main issues with *K*-means

- Poor performance in real-world scenarios due to the distance-from-cluster idea used to determine cluster membership
- Given well separable data, *K*-means finds suitable clusters




Main issues with *K*-means

- Poor performance in real-world scenarios due to the distance-from-cluster idea used to determine cluster membership
- Given well separable data, *K*-means finds suitable clusters
 - What happens when there is some overlap?



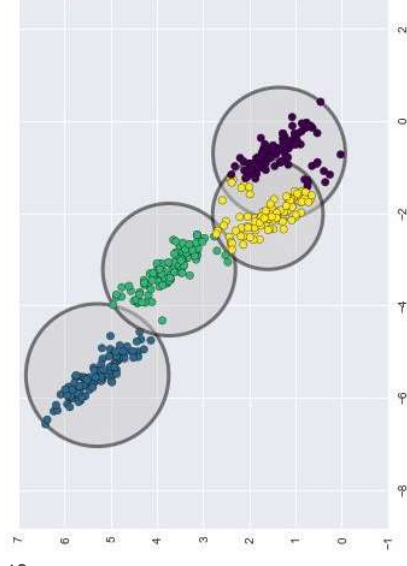
Main issues with *K*-means

- Poor performance in real-world scenarios due to the distance-from-cluster idea used to determine cluster membership
- Given well separable data, *K*-means finds suitable clusters
 - What happens when there is some overlap?
 - No measure of uncertainty cluster assignment



Main issues with *K*-means

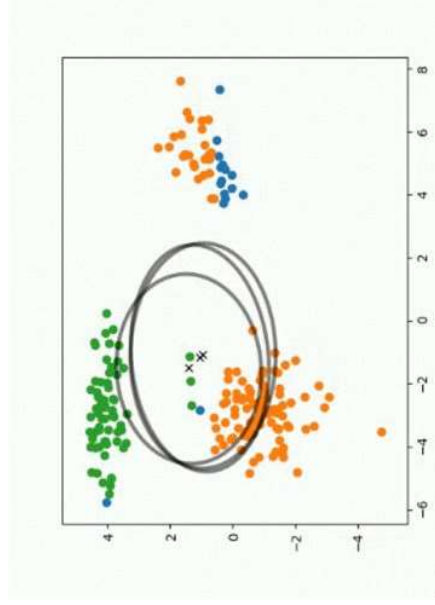
- Poor performance in real-world scenarios due to the distance-from-cluster idea used to determine cluster membership
- Given well separable data, *K*-means finds suitable clusters
 - What happens when there is some overlap?
 - No measure of uncertainty cluster assignment
- The clusters must be **circular**
 - ... and they may not be a good fit to the data
 - Not flexible enough





Gaussian Mixture Models (GMMs)

A Gaussian mixture model (GMM) attempts to find a **mixture** of multi-dimensional Gaussian probability distributions that best model any input dataset.



```
from sklearn.mixture import GMM
gmm = GMM(n_components=4).fit(X)
labels = gmm.predict(X)
```



Gaussian Mixture Models (GMMs)

A Gaussian mixture model (GMM) attempts to find a **mixture** of multi-dimensional Gaussian probability **distributions** that best model any input dataset.

Each point belongs to the cluster with a given probability.

```
probs = gmm.predict_proba(X)
print(probs[:5].round(3))
```

[0.	0.	0.475	0.525]
[0.	1.	0.	0.]
[0.	1.	0.	0.]
[0.	0.	0.	1.]
[0.	1.	0.	0.]

```
from sklearn.mixture import GMM
gmm = GMM(n_components=4).fit(X)
labels = gmm.predict(X)
```

The position and the shape of each cluster are defined by mean and covariance.





GMMs in short

Based on an **Expectation-Maximization** approach.

1. Choose starting guesses for random assignments and shapes
2. Repeat until convergence:
 - a. **E-step**: for each point, find the weights encoding the probability of membership in each cluster
 - b. **M-step**: for each cluster, update its location, normalization, and shape based on *all* data points, making use of the weights
- The resulting clusters are associated with a **smooth** Gaussian model, rather than a hard-edged sphere.
- The optimal solution may be missed, so **multiple random initializations** are used.
- No distance measures are used. The points are assigned to the clusters based on probability distributions.

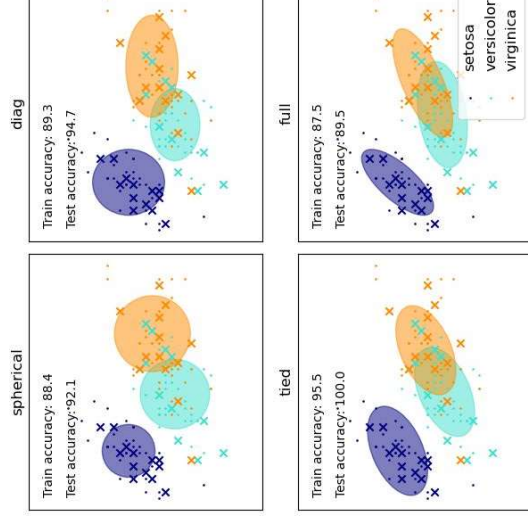


Example of resulting clusters with GMMs



Relevant choices

1. **Initialization** -- to define the initial center of the model components
 - a. k-means (can be computationally heavy)
 - b. k-means++: pick first center at random, then subsequent centers are the most distant ones
 - c. random
2. **Number of components**
3. **Covariance type** (constraint on the estimated classes)
 - a. spherical
 - b. diagonal
 - c. tied
 - d. full covariance




Exercise 2: GMMs

1. Build your first GMM to replicate the results obtained with K-means on the IRIS dataset
 - o Suggestion: `n_components` is comparable to `K`
2. Try out the different **initializations**: how do the results differ?
 - o Take a look [here](#) to see how to plot the obtained clusters.
3. Pick the best number of components by applying grid search. Use the **negative BIC** (Bayes Information Criterion) score as scoring method. Select the number of components having the lowest BIC.
 - o In sklearn: [GridSearchCV](#)
 - o Suggestion: define the function computing the BIC score and pass it to the `GridSearch`
1. How does the choice of covariance type affect the results? How do the built clusters compare with the actual dataset?
 - o Take a look at [this tutorial](#)

```
def gmm_bic_score(estimator, X):  
    """Callable to pass to GridSearchCV that will use the BIC score."""  
    # Make it negative since GridSearchCV expects a score to maximize  
    return -estimator.bic(X)
```


Exercise 3: Handling not well separable data

1. Setup the Moons dataset from sklearn as:
You should see something like



```
from sklearn.datasets import make_moons
Xmoon, ymoon = make_moons(200, noise=.05, random_state=0)
plt.scatter(Xmoon[:, 0], Xmoon[:, 1]);
```

1. Using the function `plot_gmm` from the next slide, see what happens when moving from `n_components=2` to `n_components >= 16`
 - o Is the GMM able to model the overall distribution of the input data?

Useful plotting functions

```
from matplotlib.patches import Ellipse

def draw_ellipse(position, covariance, ax=None, **kwargs):
    """Draw an ellipse with a given position and covariance"""
    ax = ax or plt.gca()

    # Convert covariance to principal axes
    if covariance.shape == (2, 2):
        U, s, Vt = np.linalg.svd(covariance)
        angle = np.degrees(np.arctan2(U[1, 0], U[0, 0]))
        width, height = 2 * np.sqrt(s)
    else:
        angle = 0
        width, height = 2 * np.sqrt(covariance)

    # Draw the Ellipse
    for nsig in range(1, 4):
        ax.add_patch(Ellipse(position, nsig * width, nsig * height,
                              angle, **kwargs))
```

```
def plot_gmm(gmm, X, label=True,
             ax=None):
    ax = ax or plt.gca()
    labels = gmm.fit(X).predict(X)
    if label:
        ax.scatter(X[:, 0], X[:, 1],
                   c=labels, s=40, cmap='viridis', zorder=2)
    else:
        ax.scatter(X[:, 0], X[:, 1],
                   s=40, zorder=2)
    ax.axis('equal')

    w_factor = 0.2 / gmm.weights_.max()
    for pos, covar, w in zip(gmm.means_,
                             gmm.covars_, gmm.weights_):
        draw_ellipse(pos, covar, alpha=w
                    * w_factor)
```

Exercise 4: GMMs for generating new data

A trained GMM describes the distribution of the input data.

1. Load the MNIST dataset (`load_digits`) and plot the first 30 samples (will serve as reference later)
2. GMMs might have troubles converging in high dimensional spaces. Apply PCA and preserve 99% of the variance.
 - a. How many principal components do you need to keep 99% of the variance?
3. Define and fit your GMM model:
 - a. Use BIC to select the best number of components to fit the reduced data. Try a few values in the range `[10, 250]`.
 - b. Fit the GMM model with the best number of components to the reduced data. Check its convergence with `gmm.converged_`
4. It's time to generate new data following the learned distribution!
 - a. Call the `sample(n_new_samples)` method on your GMM. The output should have dimensions `(n_new_samples, n_PCs)`
 - b. Apply the inverse transform the PCA to the obtained data to return to the original space
 - c. Plot the obtained digits: how similar are they to the original ones?



Some useful references

<https://jakevdp.github.io/PythonDataScienceHandbook/05.12-gaussian-mixtures.html>

<https://scikit-learn.org/stable/modules/mixture.html>

https://scikit-learn.org/stable/auto_examples/mixture/plot_gmm_pdf.html#sphx-glr-auto-examples-mixture-plot-gmm-pdf-py

<https://towardsdatascience.com/implement-expectation-maximization-em-algorithm-in-python-from-scratch-f1278d1b9137>